

Linear and K-nn models

Alessio Micheli

micheli@di.unipi.it



Dipartimento di Informatica
Università di Pisa - Italy

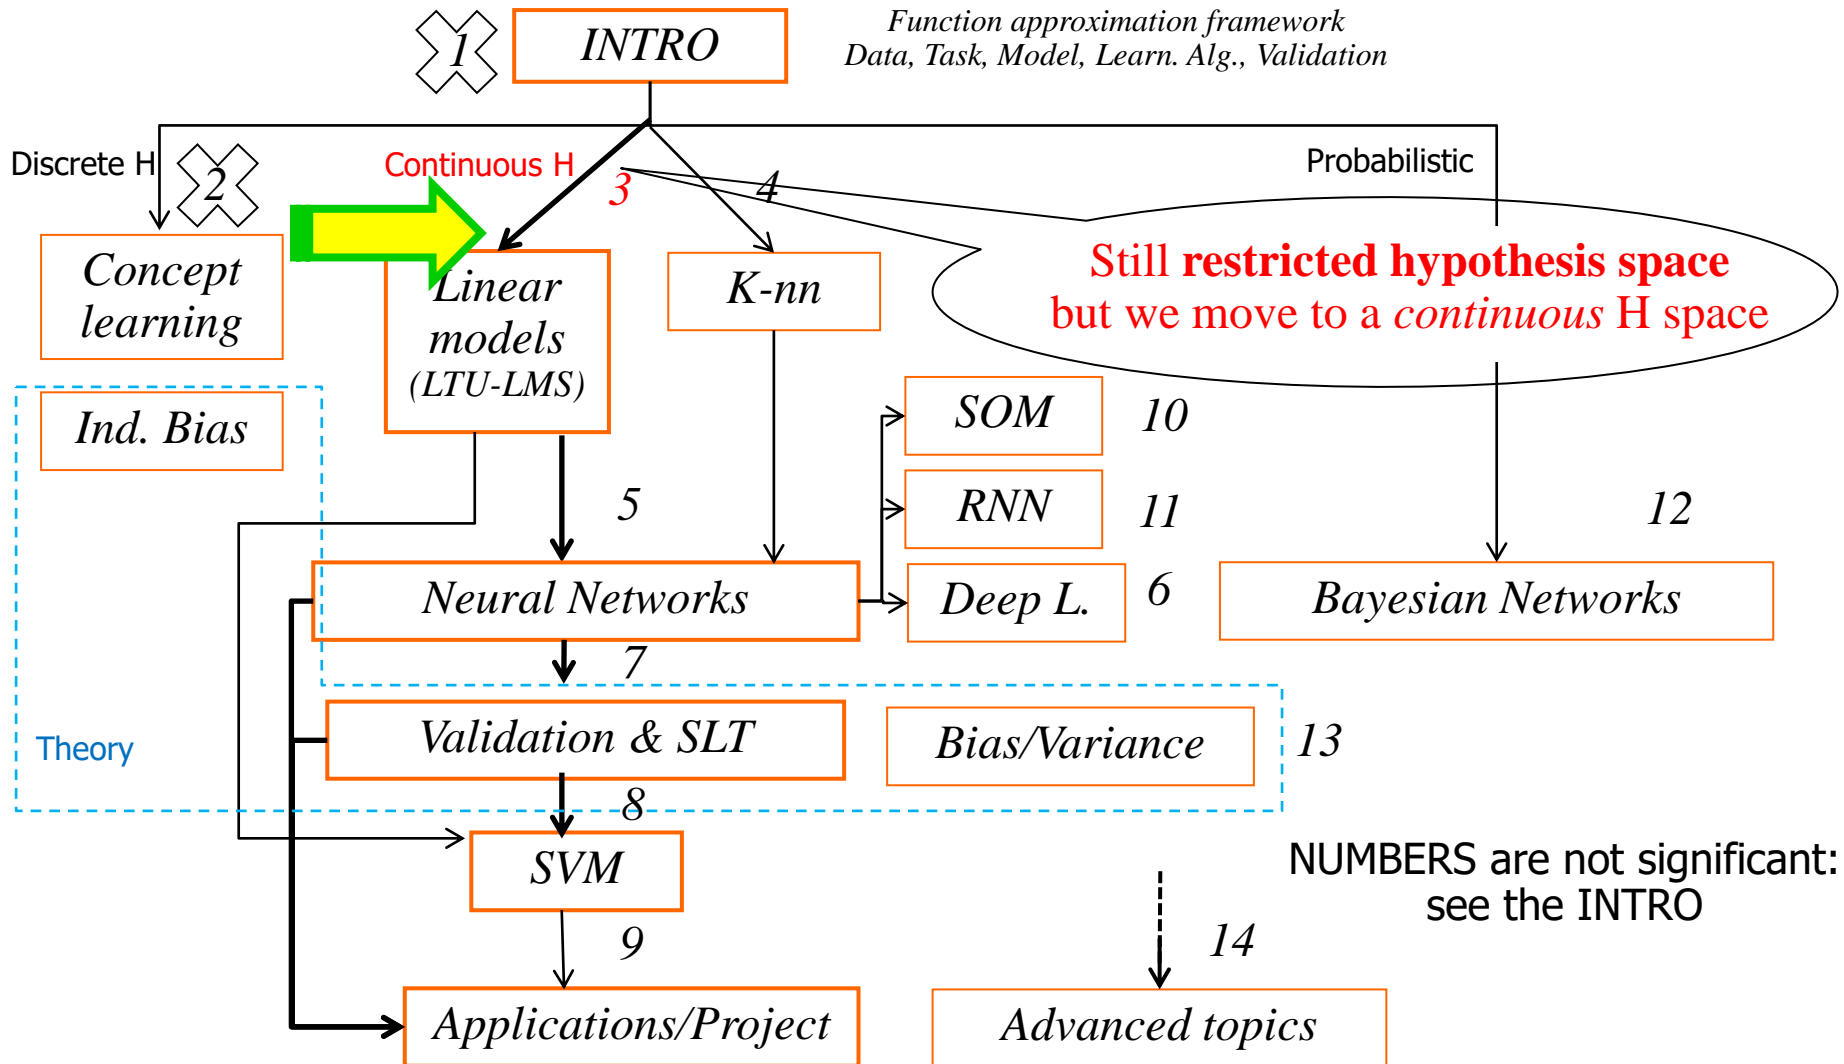
**Computational Intelligence &
Machine Learning Group**

ML Course structure

Where we go




Dip. Informatica
University of Pisa



Tasks: Supervised Learning

- **Given:** Training examples as $\langle input, output \rangle = (x, d)$ (*labeled examples*)
for an unknown function f (known only at the given points of example)
 - Target value: desiderate value d or t or y ... is given by the teacher according to $f(x)$.
- **Find:** A *good* approximation to f (a hypothesis h that can be used for prediction on unseen data x' , i.e. that is able to generalize)



- Target d (or t or y): a categorical or numerical *label*
 - **Classification:** discrete value outputs:
 $f(x) \in \{1, 2, \dots, K\}$ *classes (discrete-valued function)*
 - **Regression:** real continuous output values (approximate a real-valued target function)

Both as a *task of function approximation*

A premise on **DATA** notation

| Pattern | x_1 | x_2 | x_i | x_n |
|---------|-----------|-----------|-----------|-----------|
| Pat 1 | $x_{1,1}$ | $x_{1,2}$ | | $x_{1,n}$ |
| ... | | | | |
| Pat p | $x_{p,1}$ | $x_{p,2}$ | $x_{p,i}$ | $x_{p,n}$ |
| ... | | | | |

X is a matrix $l \times n$

l rows, n columns

$p=1..l, \quad i=1..n$

We often need to omit some indices when the context is clear, e.g.:

- Each row, generic x (vector - bold), a row in the table: (input) example, pattern, instance, sample ,..., input vector, ...
- x_i or x_j (scalar): component i or j (given a pattern x , i.e. omitting p)
- x_p or x_i (vector – bold) p -th or i -th row in the table = pattern p or i
- $x_{p,i}$ (scalar) also as $(x_p)_i$: component i of the pattern p
or also, often, $x_{p,j}$ for the component j , etc.
- For the target y we will typically use just y_p with $p=1..l$ (the same for d or t)

Linear models

- **Regression**
- **Classification**

Linear models

The linear model has been the mainstay of statistics.

- *"Despite the great inroads made by modern nonparametric regression techniques, linear models remain important, and so we need to understand them well". (Hastie)*

Plenty of studies and in many books (mathematics, statistics, numerical analysis, applicative fields, ML, ...)

- We start with the simplest form, linear in the input variables
- A baseline for learning (first: is it a linear problem?)

Regression


- **Just to see how formulate the learning problem as a **LMS** on the R_{emp}**
 - Then (next lecture) we will consider also the control of complexity
- **We formulate a first derivation in a simplified setting (univariate case)**

Repetita: regression: example

- Process of estimating of a real-value function on the basis of finite set of noisy samples
 - known pairs $(x, f(x) + \text{random noise})$

Task (exercise): find f for the data in the following table:


| x | target |
|-----|--------|
| 1 | 2.1 |
| 2 | 3.9 |
| 3 | 6.1 |
| 4 | 8.4 |
| 5 | 9.8 |
| ... | ... |




Now, we want to solve it (how to find w) in a «systematic» way

Univariate Linear Regression

- Univariate case, simple linear regression :
- We start with 1 input variable x , 1 output variable y
- We assume a model $h_w(x)$ expressed as $out = w_1 x + w_0$
- *where w are real-valued coefficients/free parameters (weights)*



- **Fitting** the data by a “straight line”



- Infinite hp space (continuous w values) but we have nice solution from classical math (going back to Gauss/Legendre ~1795!)
 - Surprisingly we can “learn” by this basic tool
 - Although simple it includes many relevant concept of modern ML and it is a basis of evolved methods in the field

Build it: Learning via LMS (I)

- **Learn** → find \mathbf{w} such that minimize **error**/empirical **loss** (best data fitting – on the training set with l examples)
- **Given** a set of l training examples $(\mathbf{x}_p, y_p) \quad p=1..l$
- **Find:** $h_{\mathbf{w}}(x)$ in the form $w_1 x + w_0$ (hence the values of \mathbf{w}) that minimizes the expected loss on the training data.
- For the loss we use the square of errors:
- Least (Mean) Square: Find \mathbf{w} to *minimize* the residual sum of squares $[\operatorname{argmin}_{\mathbf{w}} \operatorname{Error}(\mathbf{w}) \text{ in } L_2]$:

$$\operatorname{Loss}(h_{\mathbf{w}}) = E(\mathbf{w}) = \sum_{p=1}^l (y_p - h_{\mathbf{w}}(x_p))^2 = \sum_{p=1}^l (y_p - (w_1 x_p + w_0))^2$$

where x_p is p -th input/pattern/example, y_p the output for p , \mathbf{w} free par., l num. of examples


Note: to have the mean divide by l

On the notation: Indeed for the univariate case, with 1 variable: $x_p = x_{p,1} = (\mathbf{x}_p)_1$

Build it: Learning via LMS (II)

Why LMS to find the best h ?

- Least (Mean) Square: Find w to minimize the **residual sum of squares** [$\operatorname{argmin}_w \operatorname{Error}(w)$ in L_2]:



$$h_w(x)$$

$$y = w_1 x + w_0 + \text{noise}$$

Different **blue** lines will have different **green** bars.

Minimizing the **green** bars (residuals / errors)

is a way to find the best approximation/fitting of the data (i.e. our $h_w(x)$ or **blue** line).

The squares of errors $E(w)$ quantify such **green** bars:

$$E(w) = \sum_{p=1}^l (y_p - h_w(x_p))^2$$

- The method of **least squares** is a standard approach to the approximate solution of over-determined systems, i.e., sets of equations in which there are more equations than unknowns.

How to solve?


- Remember: local minimum as stationary point: the gradient is zero

$$\frac{\partial E(\mathbf{w})}{\partial w_i} = 0, \quad i = 1, \dots, \dim_input + 1 = 1, \dots, n + 1$$

- For the simple Lin. Regr. (2 free parameters)

Search the \mathbf{w} such that

$$\frac{\partial E(\mathbf{w})}{\partial w_0} = 0 \quad \frac{\partial E(\mathbf{w})}{\partial w_1} = 0$$



Convex loss function \rightarrow we have the following solution (no local minima)
(just to know that it exists!)

$$w_1 = \frac{\sum_{p:1 \rightarrow l} x_p y_p - \frac{1}{l} \sum x_p \sum y_p}{\sum_{p:1 \rightarrow l} x_p^2 - \frac{1}{l} (\sum x_p)^2} = \frac{\text{Cov}[x, y]}{\text{Var}[x]},$$

$$w_0 = \bar{y} - w_1 \bar{x}$$

$$\frac{1}{l} \sum_{p \rightarrow l} y_p \quad \frac{1}{l} \sum_{p \rightarrow l} x_p$$

Exercise: compute w_0 and w_1 according to the next slide results for the gradient (extended to l patterns)

Compute the gradient for 1 (each) pattern p



Dip. Informatica
University of Pisa

Redo this by yourself as an *Exercise*

Basic rules:

$$\frac{\partial}{\partial w} k = 0, \frac{\partial}{\partial w} w = 1, \frac{\partial}{\partial w} w^2 = 2w$$

$$\frac{\partial (f(w))^2}{\partial w} = 2f(w) \frac{\partial (f(w))}{\partial w}$$

Der. sum = sum of der.

We will call $(y - h_w(x))$
"delta"

$$\frac{\partial E(\mathbf{w})}{\partial w_i} = \frac{\partial (y - h_w(x))^2}{\partial w_i} =$$

$$= 2(y - h_w(x)) \frac{\partial (y - h_w(x))}{\partial w_i} = 2(y - h_w(x)) \frac{\partial (y - (w_1 x + w_0))}{\partial w_i}$$

$$\frac{\partial E(\mathbf{w})}{\partial w_0} = -2(y - h_w(x))$$

$$\frac{\partial E(\mathbf{w})}{\partial w_1} = -2(y - h_w(x)) \cdot x$$

Then we will sum up for l patterns $(x_p, y_p) \dots$

And we will extend to multidimensional x and w ($n=1$ here) ...

Linear model: notation for multidimensional inputs

- Assuming column vector for \mathbf{x} and \mathbf{w} (in bold)
 - Number of data l , dimension of input vector n , y_p (targets) $p=1..l$
- $$\mathbf{w}^T \mathbf{x} + w_0 = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n = w_0 + \sum_{i=1}^n w_i x_i \quad (\text{eq. 1})$$
- Note that sometimes (in NN) the transpose notation T in \mathbf{w}^T is omitted
 - w_0 is the *intercept, threshold, bias, offset*....

Often it is convenient to include the constant $x_0 = 1$ so that we can write eq.1 as :

$$\mathbf{w}^T \mathbf{x} = \boxed{\mathbf{x}^T \mathbf{w}}$$

Inner product

$$\mathbf{x}^T = [1, x_1, x_2, \dots, x_n]$$

$$\mathbf{w}^T = [w_0, w_1, w_2, \dots, w_n]$$

So, the "linear" model can be written
as a function that for each \mathbf{x}_p compute:

$$h(\mathbf{x}_p) = \mathbf{x}_p^T \mathbf{w} = \sum_{i=0}^n x_{p,i} w_i$$

Note: \mathbf{w} continuous (free) parameters
"weights"

Learning algorithm: just wait!



Dip. Informatica
University of Pisa

- For the learning algorithm in the multidimensional case for linear regression: please wait.
- We will provide it along with the learning algorithm of the linear classifier in the next few slides.

Summing up:

- Given the data set and the linear model, we can state the learning problem as **LMS problem**
- Once we find the best w parameters values, we have our $h_w(x)$ for regression purposes
- For students that need a soft intro to just LMS Regression:
 - <https://svivek.com/teaching/machine-learning/lectures/slides/linear-models/lms-regression.pdf>


Classification

What we are looking at

1. The classification by **hyperplanes**
2. See how the model works after training (“**use it**” slides)
3. How to state/formulate a **(regression)/classification learning problem** for a linear model by LMS
4. How to **derive the learning algorithm**
5. Proposing **two learning algorithms** to build a linear classifier

Problem: example

Raw Data with a Binary Response



200 points generated in \mathbb{R}^2 from an unknown distribution; 100 in each of two classes.

Can we build a rule to predict the color of future points?

Data may be generated by gaussian distribution (for each class) with different means
or by a mixture of different low variance gaussian distributions.

We reuse the linear model


- The same models (used for regression) can be used for **classification**: categorical targets (y or d), e.g. **0/1** or **-1/+1**.
- In this case we use a **hyperplane** (wx) assuming negative or positive values
- We exploit such models to decide if a point x belong to positive or negative zone of the hyperplane (to classify it)
- So we want to set w (by learning) s.t. we get good classification accuracy

Geometrical view: **hyperplane**

$w^T x$ define an hyperplane.
E.g. see the picture for 2 variables

$$x^T w = w^T x = w_0 + w_1 x_1 + w_2 x_2 = 0$$

This defines the decision boundary



Can be used to classify:

Examples $\langle (x_1, x_2), y \rangle$:


$\langle (1.0, 1.0), 1 \rangle$

$\langle (0.5, 3.0), 1 \rangle$

$\langle (2.0, 2.0), 0 \rangle$

Geometrical view: classifier

Introducing a threshold function



Using \mathbf{x}_p and
including w_0 in \mathbf{w}

In this slide \mathbf{w} is omitted from $h_{\mathbf{w}}$ (we use just h)

Examples:

$\langle (1.0, 1.0), 1 \rangle$

$\langle (0.5, 3.0), 1 \rangle$

$\langle (2.0, 2.0), 0 \rangle$

$[0,1]$ output range

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}\mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

or

$[-1,+1]$ output range

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}\mathbf{x} + w_0)$$

$$h(\mathbf{x}_p) = \text{sign}(\mathbf{x}_p^T \mathbf{w}) = \text{sign}\left(\sum_{i=0}^n x_{p,i} w_i\right)$$

Classification by linear decision boundary [repetita]




Dip. Informatica
University of Pisa

The classification may be viewed as the allocation of the input space in decision regions (e.g. **0/1**)

Example: linear separator on

2-dim instance space $\mathbf{x}=(x_1, x_2)$ in \mathbb{R}^2 , $f(\mathbf{x})=0/1$ (or $-1/+1$)



Point belonging to class 1

Separating (hyper)plane : \mathbf{x} s.t.

$$\mathbf{w}^T \mathbf{x} + w_0 = w_1 x_1 + w_2 x_2 + w_0 = 0$$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}^T \mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

[0,1]
output range

or

$$h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + w_0)$$

[-1,+1]
output range

Linear threshold unit (LTU)

Indicator functions

How many? (H): set of dichotomies induced by hyperplanes

Threshold (bias w_0)

Note that, given the bias w_0 , in the LTU

saying $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0 \geq 0$

is equivalent to say $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \geq -w_0$

with $-w_0$ as the «threshold» value

- The two forms identify the same positive zone of the classifier
- The second one emphasizes the role of the bias as a threshold value to “activate” the +1 output of the classifier.

Use it: Example (AIMA)

Classify a new data (x_1, x_2)


Find h s.t. given (x_1, x_2) return 0/-1 for *Earthquakes* and 1 for *Nuclear Explosion*

Some alg. finds this
decision boundary:

$$w_0 + w_1x_1 + w_2x_2 = 0$$



$$-4.9 + 1.7x_1 - 1x_2 = 0$$



hence $h(6, 3) = \text{sign}(w_0 + w_1 * 6 + w_2 * 3) =$
 $= \text{sign}(-4.9 + 1.7 * 6 - 1 * 3)$
 $= \text{sign}(2.3) = +1 \rightarrow \text{nuclear expl.}$



Seismic data.

x_1 body wave magnitude, x_2 surface wave magnitude

Earthquakes (white) Nuclear Explosion (black)

1982-1990 Asia

Use it: Example (Spam)

- Find $h(mail) +1$ for *spam*, -1 *not-spam*
 - Features $\Phi(mail) =$ words $[0/1]$ or phrases ("free money") $[0/1]$ or length [integer]
 - e.g. $\phi_k(\mathbf{x}) = \text{contain}(\text{word}_k)$ [bag of words representation]
- $w \rightarrow$ weight contribution of the input features to prediction
 - e.g. positive weight for "free money", negative for ".edu" or "unipi"
- $\mathbf{x}^T \mathbf{w}$ is the weight combination
- $h_w(x)$ provides the threshold to decide spam/not spam

$$h_w(\mathbf{x}) = \text{sign}\left(\sum_k w_k \phi_k(\mathbf{x})\right) > 0 \rightarrow +1 = \text{Spam} !$$




Some useful properties

(we will use them in future lectures)

$$\mathbf{w}^T \mathbf{x} + w_0 = w_1 x_1 + w_2 x_2 + w_0 = 0$$

Exercise: Draw $x_2 = -x_1 w_1/w_2 - w_0/w_2$ with different w values



A linearly separable problem

- If $w_0=0$ the line goes through the *origin* of the coordinate system.
- If $n > 2 \rightarrow$ *hyperplane* (decision boundary)
- *Scaling freedom*: the same decision boundary multiplying w by K
- w is a vector *orthogonal* to the hyperplane:

Given x_a, x_b (belonging to the sep. hyperplane):
 $\mathbf{w}^T \mathbf{x}_a + w_0 = 0$; $\mathbf{w}^T \mathbf{x}_b + w_0 = 0$ (take the diff.) \rightarrow
 $\mathbf{w}^T (\mathbf{x}_a - \mathbf{x}_b) = 0 \rightarrow$ *orthogonal vectors* (dot prod. 0)

- If it exist, in general, the solution is *not unique*: there are many possible hyperplanes separating these points: also many ML alg.!!!

Learning Algorithms

- We are going to introduce 2 learning algorithms for the regression and for the classification task using a linear model, both based on LMS:
 1. A direct approach based on **normal equation** solution
 2. An iterative approach based on **gradient descent**
- We start **redefining the learning problem** and the loss for them (for l data and multidimensional inputs)

The learning problem (classification tasks)



Dip. Informatica
University of Pisa

- **Given** a set of l training examples (\mathbf{x}_p, y_p) and a loss function (measure) L
 $y_p = \{0, 1\}$ or $y_p = \{-1, +1\}$
- **Find:** The weight vector \mathbf{w} that minimizes the expected loss on the training data

$$R_{emp} = \frac{1}{l} \sum_{p=1}^l L(h(\mathbf{x}_p), y_p)$$

- For classification: Using a piecewise constant (over $sign(\mathbf{w}^T \mathbf{x})$) for the loss can make this a difficult problem.
- Assume we still use the *least squares* (as for the regression case)


Approximating expected loss by a smooth function (#)

- Initially, we can make the optimization problem easier by replacing the original objective function L (0/1 loss) by **a smooth, differentiable function**. For example, consider the popular *mean squared error (MSE loss)*

✱:

Both losses satisfy the minimization of error (*),

Let us start with LMS avoid to introduce combinatorial problems



(*) Example:
 $h(\mathbf{x})=1$ if $\mathbf{w}^T\mathbf{x}>0$,
and for y (target)=1
 \rightarrow minimum err. is for $\mathbf{w}^T\mathbf{x}>0$

Hence no classif. error
minimizing either
0/1 loss
or MSE loss

Learning (a classifier) by Least Squares



Dip. Informatica
University of Pisa

- Find optimal values for \mathbf{w} (for fitting of training (TR) data) by *least squares*:
- **Given** a set of l training examples (\mathbf{x}_p, y_p) , **find** \mathbf{w} to *minimize* the residual sum of squares:

$$E(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 = \|\mathbf{y} - \mathbf{X}\mathbf{w}\|^2$$

Where \mathbf{x}_p is p -th input vector, y_p the output for p , \mathbf{w} free par., l num. of examples, n input dim.

– Min error: if $y_p=1$ then $\mathbf{x}_p^T \mathbf{w} \rightarrow 1$; if $y_p=0$ then $\mathbf{x}_p^T \mathbf{w} \rightarrow 0$

– **Note:** in $E(\mathbf{w})$ we do **not** use $h(\mathbf{x})$, as for regression, to hold a continuous differentiable loss (because $h(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x})$ for classification) !!!

- This is a quadratic function \rightarrow minimum always exists (but may be not unique) [see course of CM]
- \mathbf{X} is a matrix $l \times n$ with a row for each input vector \mathbf{x}_p
- Note: The same approach is used for a *regression problem*

Learning Algorithms

- We will introduce 2 learning algorithms for the regression and for the classification tasks using a linear model, both based on LMS

1. A direct approach based on **normal equation** solution

2. An iterative approach based on **gradient descent**



Normal equation & direct approach solution

- Differentiate $E(\mathbf{w})$ with respect to \mathbf{w} :
Blackboard or Exercise (a next slide).

Result synthesis:

- In the derivation we find that
$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j}$$

- We can get the **normal equation**
(point with gradient of E w.r.t $\mathbf{w} = 0$):


$$(\mathbf{X}^T \mathbf{X}) \mathbf{w} = \mathbf{X}^T \mathbf{y}$$

- If $\mathbf{X}^T \mathbf{X}$ is not singular the unique solution is given by

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} = \mathbf{X}^+ \mathbf{y} \quad \text{`+` Moore-Penrose pseudoinverse}$$

(also if \mathbf{X} is not invertible)

- Else the solution are infinite (satisfying the normal equation):
we can choose the *min norm* (\mathbf{w}) solution



See also
CM course

Direct approach by SVD

- The *Singular Value Decomposition (SVD)* can be used for computing the pseudoinverse of a matrix (X^+)

$$X = U \underset{\substack{\downarrow \\ \text{diagonal}}}{\Sigma} V^T \Rightarrow X^+ = V \Sigma^+ U^T$$

by replacing every nonzero entry by its reciprocal

- Moreover we can apply directly SVD to compute $w = X^+ y$ obtaining the minimal norm (on w) solution of least squares problem.
- Note:** THIS IS the learning alg. for the direct approach solution on w
- A practical tool. E.g. see “numerical recipes” in C and many numerical/statistical tools and scientific library (also in R, Octave, Matlab,...)
 - e.g. ARMADILLO: since 2011/12 C++ linear algebra library, NumPy etc.
- Many algorithms addressing the problems of efficiency and stability
- See also **CM** course

Solution (to find the normal eq.)

Make as an Exercise !!!(#)

$$\begin{aligned}
 \frac{\partial E(\mathbf{w})}{\partial w_j} &= \frac{\partial \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2}{\partial w_j} = \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \frac{\partial (y_p - \mathbf{x}_p^T \mathbf{w})}{\partial w_j} = \\
 &= \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \left(0 - \cancel{\frac{\partial (x_p)_1 w_1}{\partial w_j}} - \cancel{\frac{\partial (x_p)_2 w_2}{\partial w_j}} - \dots - \cancel{\frac{\partial (x_p)_j w_j}{\partial w_j}} - \dots \right) = \\
 &= \sum_{p=1}^l 2(y_p - \mathbf{x}_p^T \mathbf{w}) \left(-\frac{\partial (x_p)_j w_j}{\partial w_j} \right) = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) (x_p)_j
 \end{aligned}$$

Only the
component j
is not 0

Imposing this =0, we can easily obtain the *normal equation* (first by “sums”, then in matrix notations)

And we also obtained the gradient of E

rewritten as:

$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -\sum_{p=1}^l 2 \left((y_p - \mathbf{x}_p^T \mathbf{w}) \right) x_{p,j} = -2 \sum_{p=1}^l \delta_p x_{p,j}$$

(Note: In the original image, the term $(y_p - \mathbf{x}_p^T \mathbf{w})$ is circled with a dashed line, and a blue δ_p is placed above it.)

(we will use this form in the future, with Neural Networks)

Other approaches to LS

Many, for instances by an **iterative/gradient descent technique** we can search for:

- More efficient solutions (previous is cubic with dim of matrix X)
- Regularization (to reduce complexity of the model)
- Better approximation with noisy data by stopping searching before the minimum
- Other approaches that can be applied also to NON-LINEAR models !!!

Hence, we are going to present the second (more important) learning algorithm

Learning Algorithms

- We will introduce 2 learning algorithms for the regression and for the classification tasks using a linear model, both based on LMS
 1. A direct approach based on **normal equation** solution
 2. An iterative approach based on **gradient descent**

This approach will be the basis for fundamental approaches we will see later



Gradient descent

- Previous derivation suggest the line to construct an iterative algorithm based on :


$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})(\mathbf{x}_p)_j$$

Where \mathbf{x}_p is p -th input pattern,
 y_p the output for p , \mathbf{w} free par.,
 l num. of examples

→ Component j of pattern p , also $x_{p,j}$

- Gradient = ascent direction: we can move toward the minimum with a gradient descent (changing \mathbf{w} with $\Delta \mathbf{w} = -$ gradient of $E(\mathbf{w})$)


- *Local search*: it begins with an initial weight vector. We modify it iteratively to decrease up to minimize the error function (steepest descent).



A single w at step 0,1,2,3,4

Error surface for linear model with 2 weights (\mathbf{w})

Parabolic for the
 $E(\mathbf{w}) = E([w_0, w_1]^T)$
(convex quadratic function)



Our "compass" to find the minimum

Hypothesis space
with 2 parameters

w_0, w_1

The gradient vector

$$\Delta \mathbf{w} = - \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}} = \begin{bmatrix} -\frac{\partial E(\mathbf{w})}{\partial w_1} \\ -\frac{\partial E(\mathbf{w})}{\partial w_2} \\ -\frac{\partial E(\mathbf{w})}{\partial w_j} \\ \dots \\ -\frac{\partial E(\mathbf{w})}{\partial w_n} \end{bmatrix} = \begin{bmatrix} \Delta w_1 \\ \Delta w_2 \\ \Delta w_j \\ \dots \\ \Delta w_n \end{bmatrix}$$

We can work in a multi-dim space without the need to visualize it
P.S. w_0 is omitted above

Using the Delta Rule

- Hence, as iterative approach we will move using a learning rule based on a «delta» (changing) of w proportional to the (opposite) of the local gradient
- The «movements» will be made iteratively according to

$$w_{\text{new}} = w + \text{eta} * \Delta w$$

(or componentwise, i.e. for each w_j)

- that is the “**learning rule**”
- and eta (η) is the “step size” (learning rate) parameter (ruling the speed of our gradient descending)

Gradient descent algorithm

A simple algorithm:

1) Start with weight vector $\mathbf{w}_{\text{initial}}$ (small), fix η ($0 < \eta < 1$).

2) Compute $\Delta \mathbf{w} = -\text{"gradient of } E(\mathbf{w}) \text{"} = - \frac{\partial E(\mathbf{w})}{\partial \mathbf{w}}$ (or for each w_j)

3) Compute $\mathbf{w}_{\text{new}} = \mathbf{w} + \eta * \Delta \mathbf{w}$ (or for each w_j)

Repeat (2) until convergence or $E(\mathbf{w})$ is "sufficiently small"

- $\Delta \mathbf{w}/l$: *least mean squares* (dividing by l , that will be the standard case)
- Batch versions ($\Delta \mathbf{w}$ after each "epoch" of l training patterns)



η (η): step size = *learning rate*: speed/stability trade-off: can be (gradually) decreased to zero (guarantee convergence, avoiding oscillation around the min.): many variants will be introduced later

Batch/On-line

- For **batch version** the gradient is the sum over all the l patterns:

$$\frac{\partial E(\mathbf{w})}{\partial w_j} = -2 \left(\sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j} \right)$$

- provide a more “precise” evaluation of the gradient over a set of l data

And we upgrade the weights after this sum

- For the **on-line/stochastic version** we upgrade the weights with the error that is computed for each pattern

- hence, the 2nd pattern output is based on weights already updated from the 1st, and so ahead
- It makes progress with each examples it looks at: it can be the faster, but need smaller *eta*:

$$\frac{\partial E_p(\mathbf{w})}{\partial w_j} = -2(y_p - \mathbf{x}_p^T \mathbf{w}) x_{p,j} = -\Delta_p w_j$$

- We will see intermediate cases later (as **mini-batch**)



Examples

Batch algorithm

We update \mathbf{w} after (repeating) an "epoch" of l training data \rightarrow (blue)


On-line algorithm (stochastic gradient descent - *SGD*)

We update \mathbf{w} after each pattern p ($\Delta_p \mathbf{w}$ for each pattern \rightarrow (purple and green)



Paths over the error surface by Batch or On-line version

Learning curve examples



These are **learning curves**:
They show how the error decreases through gradient descent iterations

P.S. No relation with color in the previous slide!

Exercise: *1 is slow, 1 is unstable, 1 is good: which one?*

Gradient descent as Error correction delta rule


$$\Delta w_j = 2 \sum_{p=1}^l (y_p - x_p^T \mathbf{w}) x_{p,j}$$

$$\mathbf{w}_{\text{new}} = \mathbf{w} + \text{eta} * \Delta \mathbf{w}$$

Where $x_{p,j}$ is the component j of the input pattern p , y_p the output for p , \mathbf{w} free par., l num. of examples.

The constant 2 can be omitted

- This is an "error correction" rule (Widrow-Hoff or **delta rule**) that change each w_j proportionally to the error (target y - output):
 - E.g. (target y – output) = err=0 \rightarrow no correction
 - (input $_j$ >0) if err + (output is too low), positive delta \rightarrow increase $w_j \rightarrow$ increment the output \rightarrow less err
 - (input $_j$ >0) if err - (output is too high), negative delta \rightarrow decrease $w_j \rightarrow$ reduce output \rightarrow less err
 - ... [exercise: all the cases]
- We improve by learning from previous errors
"seeking and blundering we learn (Goethe)"




Delta-W as Error Correction Learning rule (II)




Dip. Informatica
University of Pisa

Inputs Weights ($w_0=0$)



If **misclassified** (because target is +1) $\rightarrow (1 - (-1)) \rightarrow$ High positive delta for w_1 and $w_3 \rightarrow$ increase them proportionally (with eta) to the *delta*, hence a positive value in this case
[error correction rule] !!!

e.g. (see figure):



Now is **correct** !!!

Exercise: compute the values for delta (hint: above) and for also eta

Gradient descent: final discussion



Dip. Informatica
University of Pisa

Gradient descent approach it is simple and effective *local search* approach to LMS solution and

- It allows us to search through an *infinite (continuous) hypothesis* space!
- It can be easily always applied for *continues* H and *differentiable loss*
- NOT ONLY to linear models !!!! (we will see for Neural Networks and deep learning models)
-
- *Efficient?* Many improvement are possible, e.g. Newton & quasi-newton methods; Conjugate Gradient, ...! → [CM course*](#)

Summarizing

- Model **trained** (on TR set) with LS (LMS) on $w\mathbf{x}$
 - by the simple gradient descent algorithm used for linear regression
- Model **used** for classification applying a threshold function, obtaining $h(\mathbf{x}) = \text{sign}(w^T \mathbf{x})$
- The error can be computed as *classification error* or number of misclassified patterns (not only by the Mean Square Error)

$$L(h(\mathbf{x}_p), d_p) = \begin{cases} 0 & \text{if } h(\mathbf{x}_p) = d_p \\ 1 & \text{otherwise} \end{cases} \quad \text{mean_err} = \frac{1}{l} \sum_{p=1}^l L(h(\mathbf{x}_p), d_p)$$

num_err

- **ACCURACY** = mean of correctly classified = $(l - \text{num_err})/l$

Coming back to the problem: Linear model solution

Linear Regression of 0/1 Response




Figure 2.1 (© HTF 2001):

A classification example in two dimensions. The classes are coded as a binary variable GREEN = 0, RED = 1 and then fit by linear regression. The line is the decision boundary defined by $\mathbf{x}^T \mathbf{w} = 0.5$. The red shaded region denotes that part of input space classified as RED, while the green region is classified as GREEN.

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x}^T \mathbf{w} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Linear threshold unit

The *decision boundary* is $\{\mathbf{x} \mid \mathbf{x}^T \mathbf{w} = 0.5\}$ is linear (and seems to make many errors on the training data). Is it true?

Good or bad approximation?

(#)



Dip. Informatica
University of Pisa

- Possible scenarios (we know the true target function!)
 - Scenario 1: The data in each class are generated from a Gaussian distribution with uncorrelated components, same variances, and different means.
 - Scenario 2: The data in each class are generated from a mixture of 10 gaussians in each class.
- For Scenario 1, the linear regression rule (by LS) is almost **optimal** (is the best one can do). The region of overlap is inevitable (due to errors in the input data).
- For Scenario 2, it is far too rigid: next models for it!

* Least squares corresponds to the maximum likelihood criterion if the experimental errors have a normal distribution

Linear model (in ML): Inductive Bias (alla Mitchell)



Dip. Informatica
University of Pisa


- **Language bias**: the H is a set of linear functions (may be very restrictive and rigid)
- **Search bias**: ordered search guided by the Least Squares minimization goal
 - For instance, we could prefer a different method to obtain a restriction on the values of parameters, achieving a different solutions with other properties (in particular to consider the generalization issue), ...

It shows that even for a “simple” model there are many possibilities.
We need a principled approach! (see theory of ML)...

Limitations: regression tasks for non linear problems



Dip. Informatica
University of Pisa




Too poor solution

Limitations (classification) (language bias)



Dip. Informatica
University of Pisa

- In geometry, two set of points in a two-dimensional plot are **linearly separable** when the two sets of points can be completely separated by a single line
- In general, two groups are *linearly separable* in n -dimensional space if they can be separated by an $(n - 1)$ -dimensional hyperplane.



- The linear decision boundary can provide exact solutions only for linearly separable sets of points

Example: Conjunctions

- We can represent conjunctions by the linear models, e.g.:
- Conjunctions** (see the example in the introduction lectures):

4 var.: $x_1 \wedge x_2 \wedge x_4 \leftrightarrow y$



- $1 \ x_1 + 1 \ x_2 + 0 \ x_3 + 1 \ x_4 \geq 2.5$

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{w}\mathbf{x} + w_0 \geq 0 \\ 0 & \text{otherwise} \end{cases}$$

In the plot:

2 var.: $x_1 \wedge x_2 \leftrightarrow y$

- $1 \ x_1 + 1 \ x_2 \geq 1.5$



\mathbf{w} can be learned to find this solution


Limitations: classification tasks



Dip. Informatica
University of Pisa









- Given **3 points**, can we always find a separation plane for every assignment of $f(\mathbf{x})$?
No, 3 aligned points with 0 in the middle and others 1; yes if they are not aligned (existence!).

Here NOT complete:
There are 2^3 cases



- Given **4 points**, can we always find a separation plane for every assignment of $f(\mathbf{x})$?
No (XOR)


we can find a labeling such that the linear classifier fails to be perfect

| | | XOR |
|---|---|------|
|  |  | 00 0 |
|  |  | 01 1 |
|  |  | 10 1 |
|  |  | 11 0 |

Exercise: try to redo it by yourself!

How to extend the linear model

- Note that in $h_{\mathbf{w}}(x) = w_1 x + w_0$ or $h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \cdot \mathbf{x}$
- As Statistical Parametric models:
"linear" does not refer to this (red) straight line, but rather to the way in which the regression coefficients \mathbf{w} occur in the regression equation



- Hence, we can use also transformed inputs, such as x, x^2, x^3, x^4, \dots with *non-linear* relationship inputs and output, holding the learning machinery (Least Square solution) developed so far...

$$h_{\mathbf{w}}(x) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

polynomial regression

A generalization (**LBE**)

(shown for regression*)



Dip. Informatica
University of Pisa

- Basis transformation: **linear basis expansion (LBE)**:

$$h_w(x) = \sum_{k=0}^K w_k \phi_k(x)$$

- Augment the input vector with additional variables which are transformations of x according to a function ϕ ($\phi_k: R^n \rightarrow R$)
- E.g.
 - Polynomial representation of x : $\phi(x) = x_j^2$ or $\phi(x) = x_j x_i$, or
 - Non-linear transformation of single inputs: $\phi(x) = \log(x_j)$, $\phi(x) = \text{root}(x_j)$,
 - Non-linear transformation of multiple input: $\phi(x) = \|x\|$
 - *Splines*, ...,
- Number parameters $K > n$ (before it was n)
- The model is *linear in the parameters (also in ϕ , not in x)*: we can use the **same learning alg.** as before!
- Note: it can be applied for regression (here) or classification (Exercise: HOW?)

Basis Expansion: examples



Dip. Informatica
University of Pisa

- Basis transformation: **linear basis expansion**:

$$h_{\mathbf{w}}(\mathbf{x}) = \sum_{k=0}^K w_k \phi_k(\mathbf{x})$$

EXAMPLES:

- [1-dim \mathbf{x}] $\phi_j(x) = x^j$.

$$h(\mathbf{x}) = w_0 + w_1x + w_2x^2 + \dots + w_Mx^M = \sum_{j=0}^M w_jx^j$$

1-dim polynomial regression ($K=M$)

Already seen in
the introduction

- Or “any other”, e.g. $\phi(\mathbf{x}) = \phi([x_1, x_2, x_3])$

$$h(\mathbf{x}) = w_1x_1 + w_2x_2 + w_3\log(x_2) + w_4\log(x_3) + w_5(x_2x_3) + w_0$$

Basis Expansion criticism

- Which ϕ (ϕ) ? Toward the so called “**dictionary**” approaches
- **PROS**: Can model more complicated relationships (than linear): it is more expressive.
- **CONS**: With *large* basis of functions, we easily risk *overfitting*, hence we require methods for controlling the complexity (*) of the model
 - *Curse of dimensionality* (the volume of the problem space increases so fast that the available data become sparse, the data became no more sufficient to support the model complexity) → we will see later
 - Φ are *fixed before* observing training data
 - versus adaptive /non-linear in parameters e.g. NN!
- *We will see the alternative NN and SVM solutions!*

(*) Whereas **complexity** is **not** for the computational cost but a measure of the flexibility of the model to fit the data (see the VC-dim)

Improvements: How to control model complexity? **Important !**



Dip. Informatica
University of Pisa

- Many approaches... e.g. coefficient shrinkage:
- **Ridge regression**: (**Tikhonov regularization**): smoothed model
→ possible to add constraints to the sum of value of $|w_j|$ penalizing models with high values of $|w|$, i.e. favoring "sparse" models using less terms due to weights $w_j = 0$ (or close to 0) (it means a less complex model)

"Error" data term [(M)SE] Regularization/penalty term

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda \|\mathbf{w}\|^2 \longrightarrow \sum w_j^2$$

Lambda (λ): regularization (hyper)parameter
(a small positive value chosen by the "model selection" phase)

The sum is over the number of w

- Note that for the **objective function** we use here the name *Loss* (used for the model training cost function) to distinguish from the *Error E* (useful to evaluate the model error and used for the *data term* inside this Loss).
Hence, differently from previous assumptions the two terms are not more equivalent in our use (in the course).

Tikhonov Regularization: Solving it



Dip. Informatica
University of Pisa

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda ||\mathbf{w}||^2$$

- For the direct approach:

$$\mathbf{w} = \boxed{(\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1}} \mathbf{X}^T \mathbf{y} \quad \text{-----} \rightarrow \text{this matrix is always invertible}$$

- For the gradient approach:
 - As an **Exercise** compute the gradient of the two terms (error and penalty terms) of the $Loss(\mathbf{w})$ w.r.t. weights w_i separately, using eta only for term E
 - Formulate the new update rule:
 - You will obtain: $\mathbf{w}_{new} = \mathbf{w} + eta * \Delta \mathbf{w} - 2 \lambda \mathbf{w}$
 - That is a **weight decay** technique (basically add $2\lambda \mathbf{w}$ to the gradient)
 - E.g. with 0 gradient, it decreases the value of each w with a fraction of the old w

Andrej Nikolaevič Tikhonov



Russian mathematician
1906 –1993

Tikhonov regularization: trade-off



Dip. Informatica
University of Pisa

$$Loss(\mathbf{w}) = \sum_{p=1}^l (y_p - \mathbf{x}_p^T \mathbf{w})^2 + \lambda ||\mathbf{w}||^2$$

- Note the balancing (trade-off) between the two terms:
 - Small **lambda** (λ) value \rightarrow minimizing the loss the focus is on obtaining a small error data term (first term, minimize just the training error) with a too complex model (high norm of the weights), the risk is of **overfitting**,
 - High **lambda** (λ) \rightarrow minimize the loss the focus is on the second term, hence the data error (first term) could grow too much, i.e. moving to **underfitting**
 - The trade-off is ruled by the value of **lambda** (λ)
 - We will see soon some examples
- The main advantage is that we have a concrete realization of the control of model complexity, easy to be implemented and of general applicability.

Tikhonov and SLT


- The penalty term **penalizes high value of the weights** and tends to drive all the weights to smaller values
 - E.g. Some weights values can go even to zero
- It **implements** a control of the model complexity
- This leads to a model with **less (or proper) VC-Dim** (with a trade-off obtained through **just a (1) parameter** that you can control: the λ)
- **Exercise:** connect Tikhonov regularization with the slide on SLT (see the introduction) to see
 - why this can help to have a better bound on R , and
 - how lambda values can rule the underfitting/overfitting cases
- **Exercise:** derive the new learning rule with weight decay using the Tikhonov loss: computing again the partial derivatives of $Loss$ with respect to w , separating data term (\bullet eta) and penalty term (\bullet λ)

Slide from Intro slides: 9th Order Polynomial



Dip. Informatica
University of Pisa

As to have Lambda 0



We have already seen it

$E(\mathbf{w}) = 0$ on training data and **overfitting**

Regularization effect

$$\ln \lambda = -18$$




Dip. Informatica
University of Pisa

- 9th Order Polynomial

Lambda small positive


$$\lambda = 0.0000000152$$



Once regularized with a suitable value of lambda
it works well! We avoid the overfitting

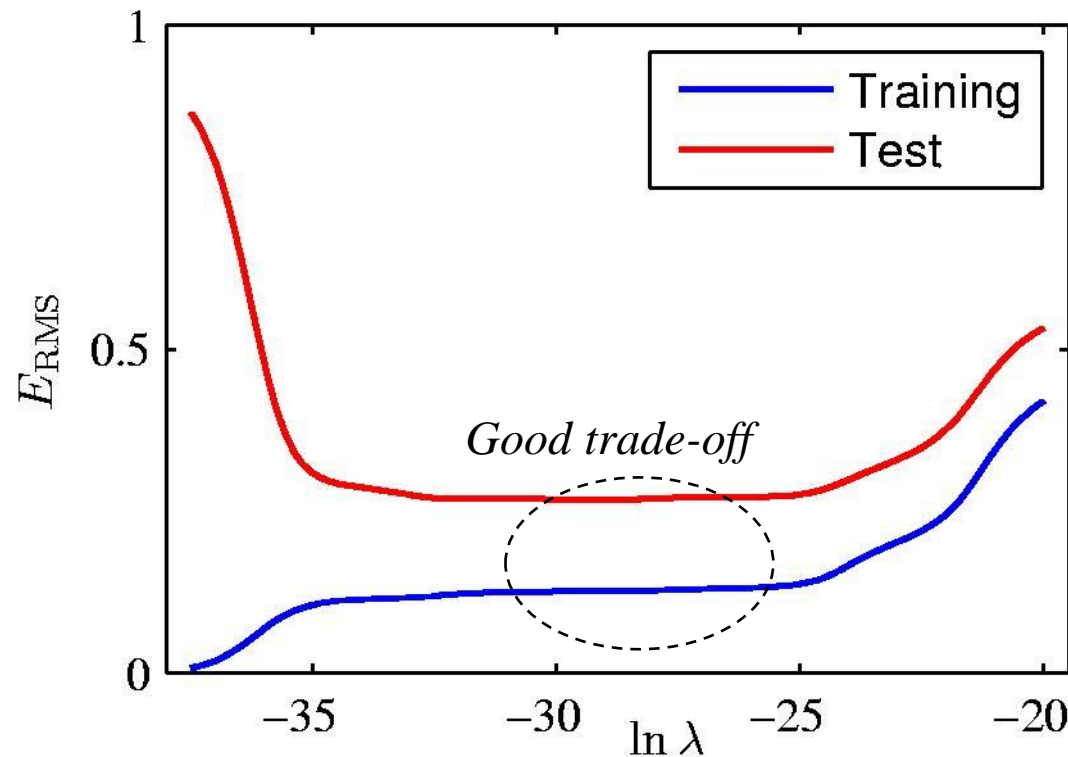
Too High lambda ($\ln \lambda = 0$)

Very high lambda, close to 1



*But we need a trade-off ...
(because a too high lambda leads as to underfitting)*

Regularization: E_{RMS} vs $\ln \lambda$



Low lambda \rightarrow overfitting

High lambda \rightarrow underfitting


Polynomial Coefficients

| | $\ln \lambda = -\infty$ | $\ln \lambda = -18$ | $\ln \lambda = 0$ |
|---------|-------------------------|---------------------|-------------------|
| w_0^* | 0.35 | 0.35 | 0.13 |
| w_1^* | 232.37 | 4.74 | -0.05 |
| w_2^* | -5321.83 | -0.77 | -0.06 |
| w_3^* | 48568.31 | -31.97 | -0.05 |
| w_4^* | -231639.30 | -3.89 | -0.03 |
| w_5^* | 640042.26 | 55.28 | -0.02 |
| w_6^* | -1061800.52 | 41.32 | -0.01 |
| w_7^* | 1042400.18 | -45.95 | -0.00 |
| w_8^* | -557682.99 | -91.53 | 0.00 |
| w_9^* | 125201.43 | 72.68 | 0.01 |

Other Regularization Tech. for Linear Models



Dip. Informatica
University of Pisa

- Ridge regression ($\| \cdot \|_2$)
 - Lasso ($\| \cdot \|_1$)
 - Elastic nets (use both $\| \cdot \|_1$ and $\| \cdot \|_2$)
- 
-
- We will introduce them later in the course (or in CM course), anyway:
 - The *L2 norm* penalizes the square value of the weight and tends to drive all the weights to smaller values.
 - On the other hand, the *L1 norm* penalizes the absolute value of the weights and tends to drive some weights to exactly zero (while allowing some weights to be large) → toward feature selection!
 - Unfortunately $\| \cdot \|_1$ (absolute value) introduce a non differentiable loss → needs other approaches



Improvements (cnt)



Dip. Informatica
University of Pisa

- Inputs with added noise (*data augmentation*)
- Derived inputs
a small number of new variables is used in place of the x inputs, which are a linear combination of x .
 - Principal Component Regression
 - Partial Least Squares

Multi-class task (#)



Dip. Informatica
University of Pisa

Two very simple approaches for multi-class:

- Class 1-of-K rep: {red, green, blue} \rightarrow (0,0,1), (0,1,0), (1,0,0). \rightarrow solve 3 linear models
- **OVA**: "one-vs-all": a discriminant function for each class

built on top of real-valued binary classifiers


- train K different binary classifiers, each one trained to distinguish the examples in a single class from the examples in all remaining classes.
- to classify a new example, the K classifiers are run, and the classifier which outputs the largest (most positive) value is chosen.
- **AVA**: "all-vs-all" = "one-versus-one":
 - each classifier separates a pair of classes. Apply it to all the pairs: $K(K-1)$
[How many training?]
 - to classify a new example, all the classifiers are run, and the winner is the one with the max sum of outputs versus all the other classes **OR** the class with most votes
 - training data set for each classifier is much smaller

Criticism (##)



Dip. Informatica
University of Pisa

The problem can become not easy



- Masking: classes can be masked by others (for high K)
- A wide array of more sophisticated approaches for multiclass classification exists ...
- Some models can deal directly with multi-output

Other learner models for classification (#)



Dip. Informatica
University of Pisa

- Linear Discriminant Analysis (also multi-class)
- Logistic regression
 - $P(y/x)$ starting from modeling the class density as a know density

Extensions: ML Models

(pro future) (#)



Dip. Informatica
University of Pisa

In ML course:

- Perceptron (Rosenblatt 1958, biological inspiration):
 - “minimize (only) misclassifications” algorithm
 - basis for Neural Networks (set of units with layers)
 - Adaptive *basis expansions* (*phi learned by training*)
 - Feature representation learning in each layer (deep learning concept)
 - Gradient descent approach for learning
- SVM (Vapnik 1996):
 - *regularization* via the concept of maximum margin: Maximize the gap (margin) between the two classes on the training data.
 - enlarge the feature space via *basis expansions* (e.g. polynomials).
- NN and SVM models realize (also) a flexible ***non-linear*** approximation for classification and regression problems

Lessons Learned

- We concretely showed that is possible to *learn* by tuning free parameters \mathbf{w} , searching over a continuous space guided by a loss f .
 - How to state/formulate a regression/classification learning problem for a linear model by LMS
 - How to derive the learning algorithm (basic univariate, direct and **gradient** based forms)
 - Two learning algorithms
 - How the model works after training
 - Linear models have limitations (strong language bias)
- How to extend to non-linear modeling (by **LBE**)
- How to implement a **regularization approach** changing the loss f .
- The role of regularization for the **control of the model complexity**

Also useful for your self assessment (try to check if you can answer)

ML Course structure

Discussion on Linear Model

In the *file rouge* of ML, 2 main concepts up to now:

- Basis expansion → (implement) more flexibility
- Regularization → (implement) control of complexity
- The cases of the *linear models* provided an *instance* routed in the classical mathematical approaches but the we will find again these concepts, by different models and implemented in different/similar forms, in the modern ML (e.g. for NN & SVM in our course)!
- Now we move on the other extreme, toward a very flexible approach

K-nn

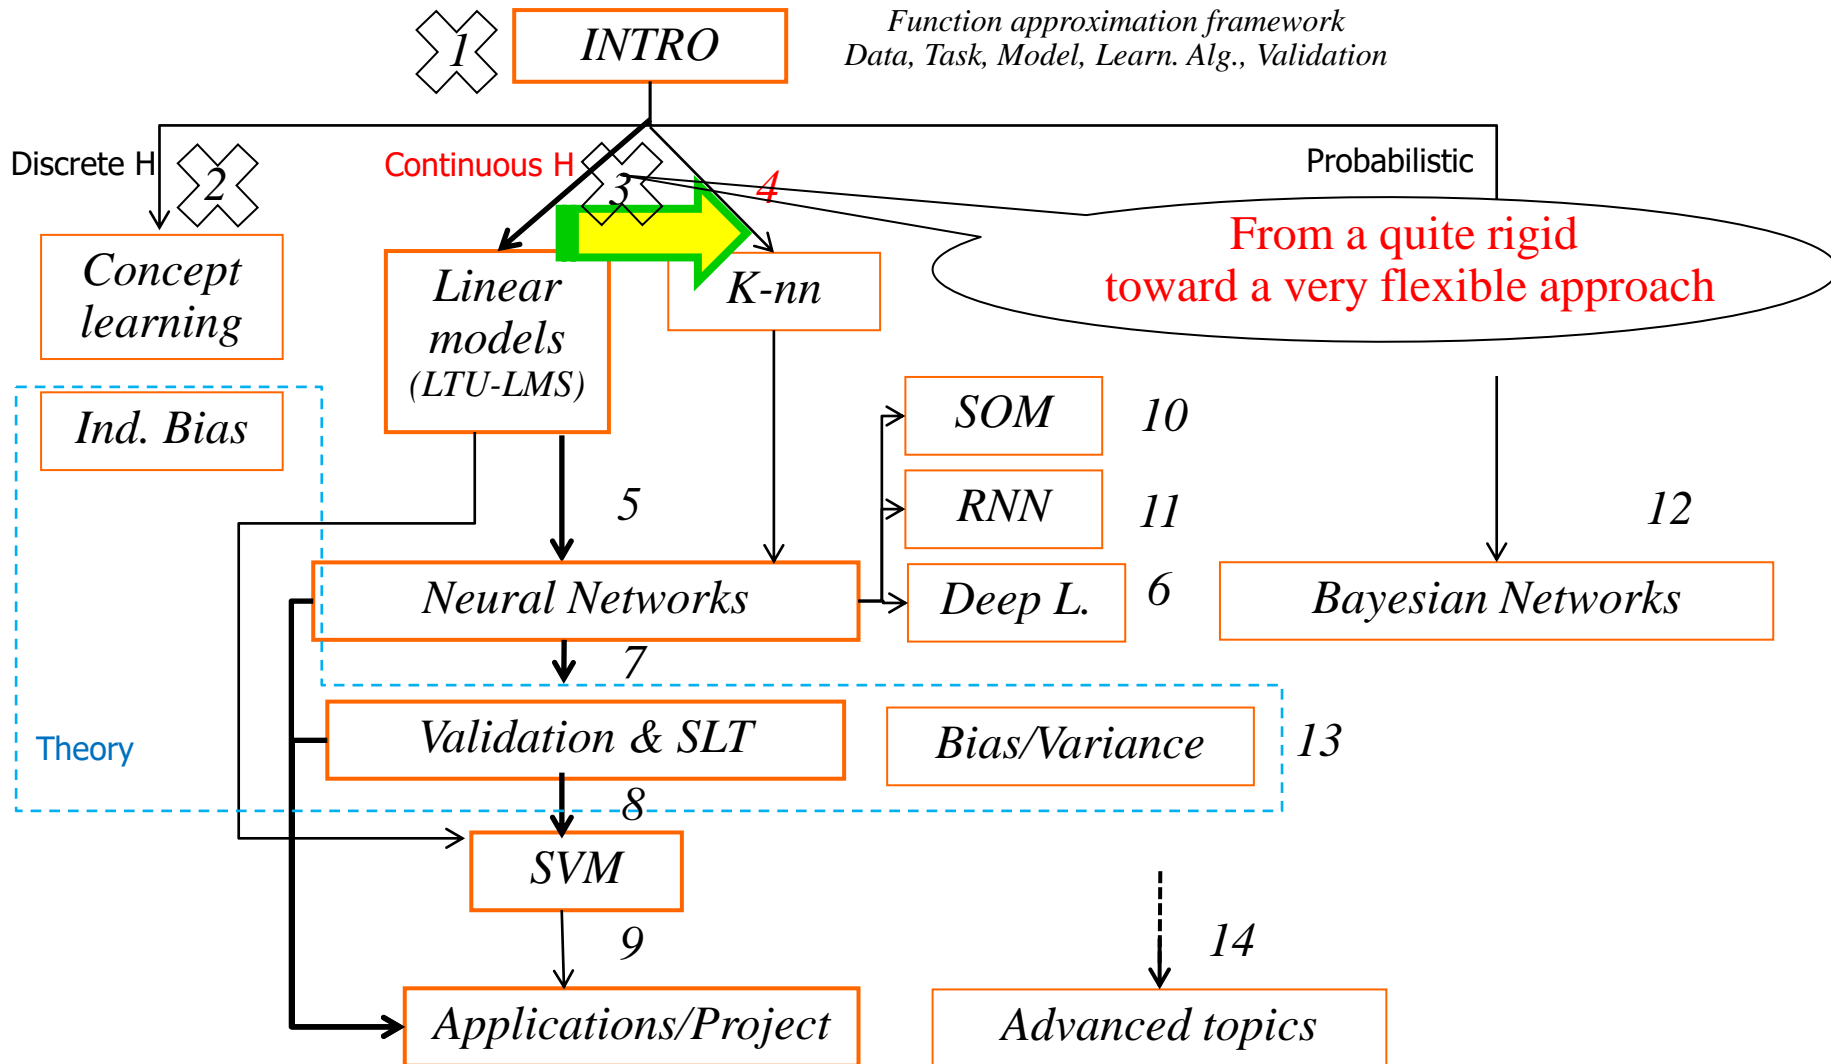
Still simple but flexible (and local)

ML Course structure

Where we go




Dip. Informatica
University of Pisa



Repetita: Problem: example

Raw Data with a Binary Response



200 points generated in \mathbb{R}^2 from an unknown distribution; 100 in each of two classes.

Can we build a rule to predict the color of future points?

Data may be generated by gaussian distribution (for each class) with different means
or by a mixture of different low variance gaussian distributions.

Repetita: Find by LS a line to separate

Linear Regression of 0/1 Response




Figure 2.1 (© HTF 2001):

A classification example in two dimensions. The classes are coded as a binary variable GREEN = 0, RED = 1 and then fit by linear regression. The line is the decision boundary defined by $\mathbf{x}^T \mathbf{w} = 0.5$. The red shaded region denotes that part of input space classified as RED, while the green region is classified as GREEN.

$$h(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x}^T \mathbf{w} > 0.5 \\ 0 & \text{otherwise} \end{cases}$$

Linear threshold unit

The decision boundary is $\{\mathbf{x} \mid \mathbf{x}^T \mathbf{w} = 0.5\}$ is linear (and seems to make many errors on the training data). Is it true?

Learning... timing

- **LEARNING ALG.**
- Timing.
 - **Eager:** Analyze the training data and construct an explicit hypothesis.
 - **Lazy:** Store the training data and wait until a test data point is presented, then construct an ad hoc hypothesis to classify that one data point.

1-Nearest Neighbor

The algorithm:

- Simply store the training data $\langle \mathbf{x}_p, y_p \rangle$ $p=1 \dots l$
- Given an input \mathbf{x} (with dimension n)
- Find the nearest training example \mathbf{x}_i
 - Find i s.t. we have $\min d(\mathbf{x}, \mathbf{x}_i) \rightarrow i(\mathbf{x}) = \arg \min_p d(\mathbf{x}, \mathbf{x}_p)$
 - E.g. Euclidian distance :
- Then output y_i

$$d(\mathbf{x}, \mathbf{x}_p) = \sqrt{\sum_{t=1}^n (x_t - x_{pt})^2} = \|\mathbf{x} - \mathbf{x}_p\|$$

↓
Pattern \mathbf{x}_p , component t

1-nn

1-Nearest Neighbor Classifier




Figure 2.3:

The same classification example in two dimensions as in Figure 2.1.


GREEN = 0, RED = 1

- Very flexible !
- No misclassifications in TR data: 0 training error: what for test data ?
- Decision boundaries is *not* linear: it is quite *irregular*
- May be unnecessary noisy (e.g. for scenario 1)

1-nn vs 5-nn example

- 1-nn return + for x_q
- 5-nn return – for x_q

Smoothing over a set of
neighbors for noisy data



K-Nearest Neighbors

- A natural way to classify a new point is to have a look at its neighbors,
- and take an average:

$$avg_k(\mathbf{x}) = 1/k \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} y_i$$

- where $N_k(\mathbf{x})$ is a neighborhood of \mathbf{x} that contains exactly k neighbors (closest patterns according to *distance* d):
 - k-nearest neighborhood: **K-nn**.
- If there is a clear dominance of one of the classes in the neighborhood of an observation \mathbf{x} , then it is likely that the observation itself would belong to that class, too. Thus the classification rule is the **majority voting** among the members of $N_k(\mathbf{x})$. As before,

$$h(\mathbf{x}) = \begin{cases} 1 & \text{—if } avg_k(\mathbf{x}) > 0.5 \\ 0 & \text{_____otherwise} \end{cases} \quad \text{for targets } y \text{ in } \{0,1\}$$

- For regression task: use directly the *avg*: mean over K-nn

15-nn

15-Nearest Neighbor Classifier




Figure 2.2:


The same classification example in two dimensions as in Figure 2.1.

GREEN = 0, **RED** = 1

The predicted class is hence chosen by majority vote amongst the 15-nearest neighbors.

- Still very flexible !
- Some misclassifications in TR data
- Decision boundaries is *not* linear: it is still quite, although less, *irregular*
- Decision boundary adapts to the local densities of the classes

Voronoi diagram



Each cell consisting of all points closer to x than to any other patterns.

The segments of the Voronoi diagram are all the points in the plane that are equidistant to two patterns.

Implicitly used by K-NN



K-nn for multi-class

- Return the class most common amongst its k nearest neighbors

$$h(\mathbf{x}) = \arg \max_v \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} \mathbf{1}_{v, y_i}$$

$$\mathbf{1}_{v, y_i} = \begin{cases} 1 & \text{if } v = y_i \\ 0 & \text{otherwise} \end{cases}$$

We count the classes in the neighbor (by the $\mathbf{1}_{v, y}$ for each v)
taking the most frequent class ($\arg \max$)

K-nn variants: Weighted distance



Dip. Informatica
University of Pisa

- It can be useful to weight the contributions of the neighbors, so that the nearer neighbors contribute more to the average than the more distant ones.

$$h(\mathbf{x}) = \arg \max_v \sum_{\mathbf{x}_i \in N_k(\mathbf{x})} \mathbf{1}_{v, y_i} \cdot \frac{1}{d(\mathbf{x}, \mathbf{x}_i)^2}$$

If $d = 0$ for a i , return y_i

$$\mathbf{1}_{v, y_i} = \begin{cases} 1 & \text{if } v = y_i \\ 0 & \text{otherwise} \end{cases}$$

K-nn: An extreme

- Not a global hypothesis for all the instances → no model to be fit
 - We need to memorize the input examples
- Local estimations (by locally constant functions) vs global linear approximation/estimation of the target function (over the instance space)
- Lazy, memory based, instance-based, **distance-based methods**

K-nn versus Linear

Discussion on linear versus k-nn models

Two extremes of the ML panorama:

- Rigid (low variance) versus flexible (high variance):
 - In K-nn with small k few points can change the decision boundary
 - We may pay a price for this flexibility
- Eager versus lazy
- Parametric versus instance-based
- Linear regression uses 3 parameters to describe its fit in the example of Fig. 2.1, w_0, w_1, w_2 (or $n+1$ in general)
Does K-nn use 1 (the value of k here)?
- More realistically, K-nn uses l/k “effective number of parameters” (Hastie-Tibshirani-Friedman 2001)*
 - where l is used the number of data (*N in the book*)

LS linear vs K-nn with various k values

Figure 2.4: *Misclassification curves for the simulation example used in Fig. 2.1, 2.2 and 2.3.*

Training set of size 200

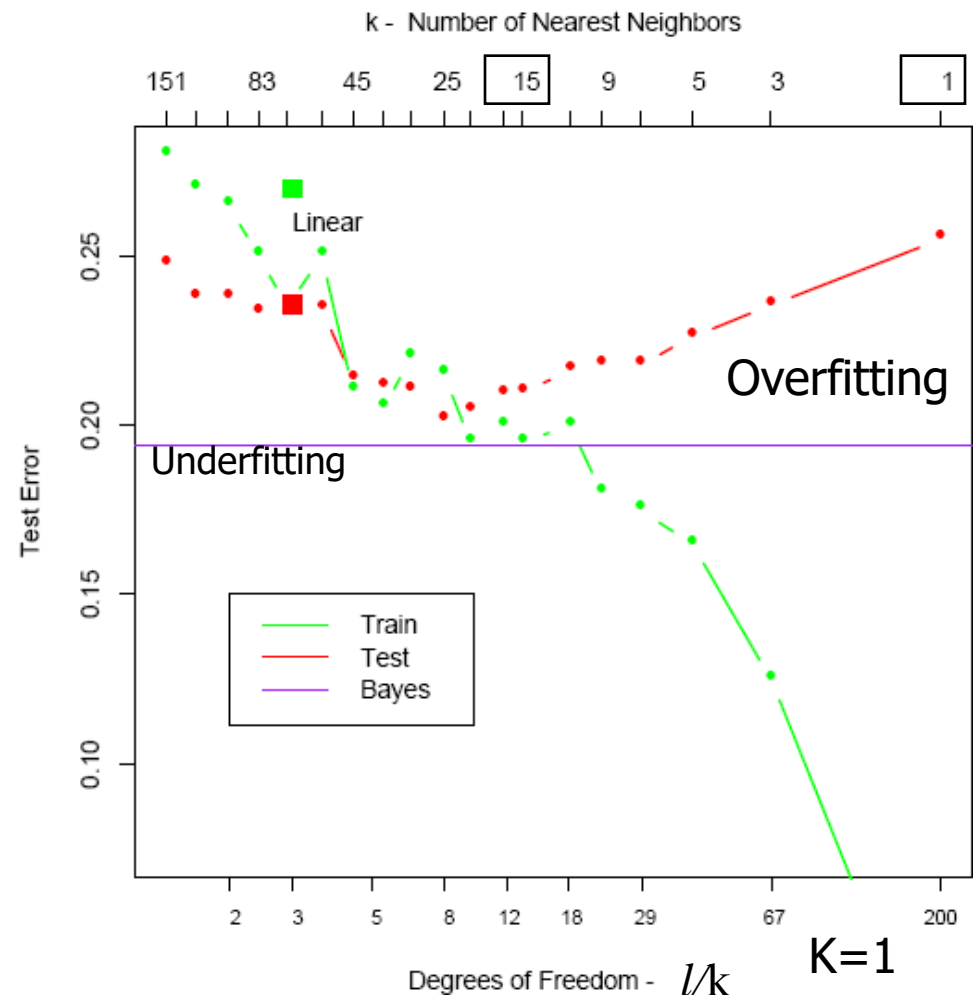
Test set of size 10.000.

The **red** curves are test and the **green** are training error for k-NN classification (changing K).

The results for linear regression are the bigger green and red dots at three degrees of freedom.

The **purple** line is the optimal Bayes Error Rate (See next slides)

Note how we move from underfitting to overfitting moving the values of k (i.e. the rate l/k): more flexibility allows to find the best result if we control "complexity" by K

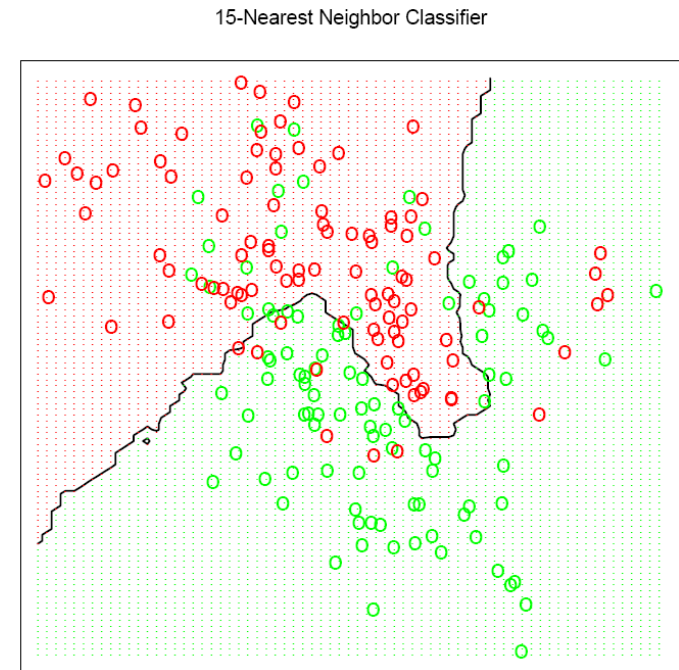



Bayes error rate

Bayes optimal solution (called Bayes classifier):

- If we know the density $P(x,y)$ we classify to the most probable class, using the conditional (discrete) distribution as:
 - Output the class v s.t. is $\max P(v/x)$ (v in $\{C_1, C_2, C_3 \dots C_K\}$)
- The error rate of the (optimal) Bayes classifier is called the Bayes rate.
- I.e. the minimum achievable error rate given the distribution of the data (assuming that the generating density is known !!!)
- Note on K-nn: we see that the k-nn classifier directly approximates this solution (a majority vote in a nearest neighborhood amounts to exactly this), except that
 - conditional probability at a point is relaxed to conditional probability within a neighborhood of a point, and
 - probabilities are estimated by training-sample proportions.

Bayes Optimal Classifier: results



- The optimal Bayes decision boundary for the simulation example.
- Since the generating density is known for each class, this boundary can be calculated exactly
- 15-nn was close to this (indeed it shown a min. test err in the plot a couple of slides ago)

Inductive Bias of k-nn


- The assumed **distance** tells us which are the most similar examples
- The classification is assumed similar to the classification of the neighbors according to the assumed metric
 - E.g. we can use other metric than the Euclidian
 - Symbolic data require “ad-hoc” metrics (e.g. Hamming distance between two strings of equal length is the number of positions for which the corresponding symbols are different).
- Next slides:
 - criticism and limitations

1. Scale changes and 2. other metrics



Dip. Informatica
University of Pisa

- **Scale changes:** Domain knowledge dependent choice
- **If** variables should contribute equally:
 - Pay attention to disparity in the ranges of each variables
 - Rescale data to equalize inputs ranges (\rightarrow change the metric)!
E.g. mean zero and variance 1 normalization
- Variable scaling can have a high impact (i.e. k-nn is fragile even with respect to basic preprocessing)



(some) Limits of K-nn: computational cost



Dip. Informatica
University of Pisa

- Note that K-nn makes the local approximation to the target function for each new example to be predicted:
 - The computational cost is deferred to the **prediction phase!**

Moreover: high retrieval cost:

- Computationally intensive (in time) for each new input: computing the distances from the test sample *to all* stored vectors
 - The time is proportional to the number of stored patterns
 - “ad-hoc” proximity search algorithms to optimize
 - E.g. by indexing the patterns
- Cost in space (all the training data)

(some) Limits of K-nn



Dip. Informatica
University of Pisa

- K-nn models offer little interpretation.
 - Subjectivity of interpretation
 - Dependence on the metric

(some) Limits of K-nn: curse of dim.

K-nn provides a good approximation if we can find a significant set of data close to any \mathbf{x} , with dense sampling

When can it fail?

- When we have a lot of input variables (high n , *i.e. high dim.*), K-nn methods often fails because of the ***curse of dimensionality***:

Many manifestations of this problem: we will examine a few

1. *It is hard to find nearby ("similar") points in high dimensions!*
2. *Low sampling density for high-dim data*
3. *Irrelevant features issue*

*These are important in terms of ML,
as they inherently affect the **generalization capability** !!!*

(some) Limits of K-nn: curse of dim.



1. It is hard to find nearby points in high dimensions!

- K-nearest neighbors can fail in high dimensions, because it becomes difficult to gather K observations “close” (in terms of of variables values, i.e. “similar”) to a query point x_q :
 - near neighborhoods tend to be **spatially large**, and estimates are not longer local
 - reducing the spatial size of the neighborhood means reducing K, and the variance of the estimate increases (\rightarrow overfitting).
- Why? See the next slide...

(some) Limits of K-nn: curse of dim.

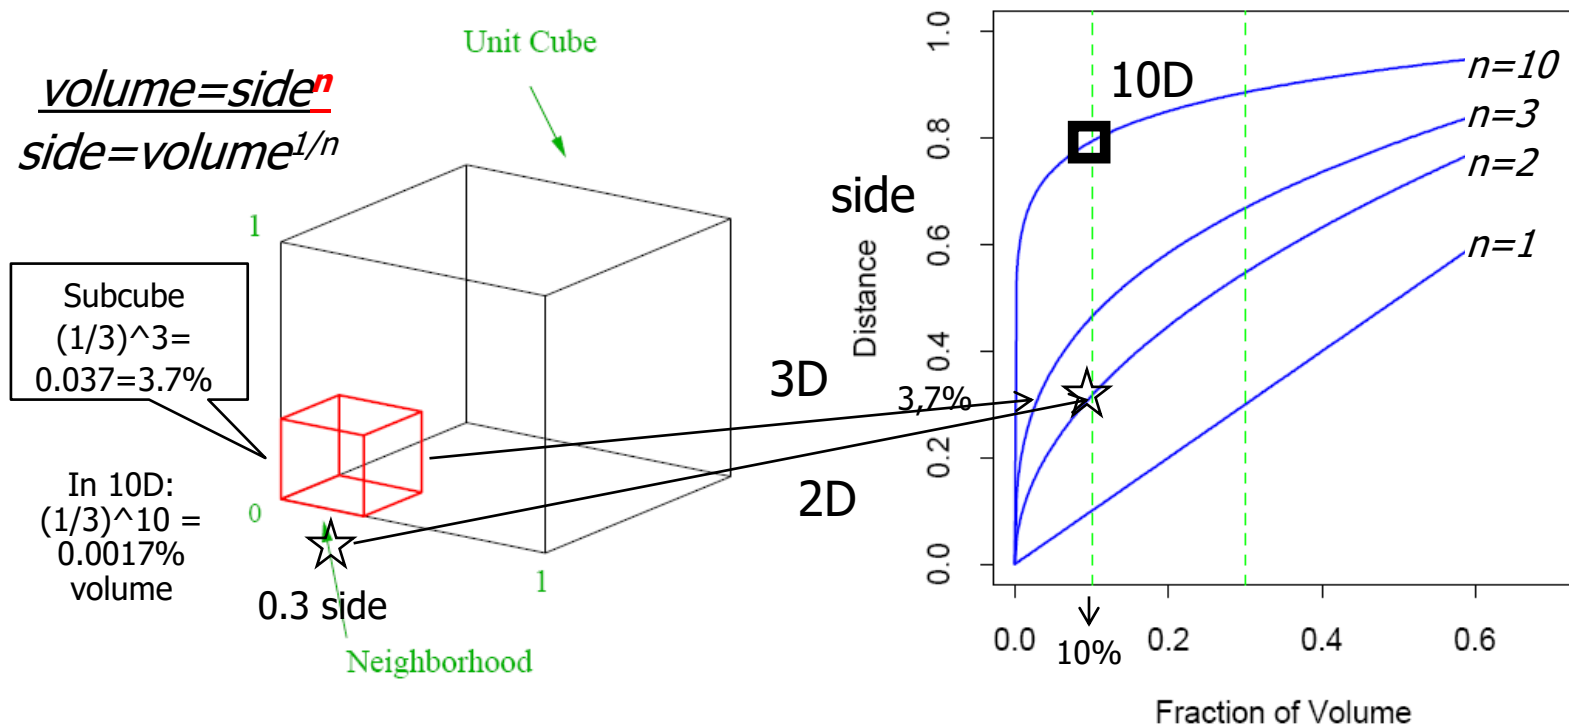


Figure 2.6: The curse of dimensionality is well illustrated by a subcubical neighborhood for uniform data in a unit cube.

The figure on the right shows the **sidelength** ($=r^{1/n}$) of the subcube needed to capture a **fraction r** of the volume of the data, for different dimensions n .


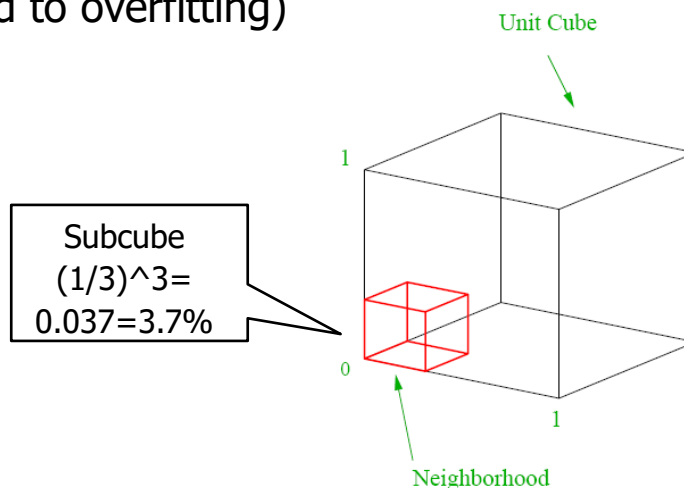
In ten dimensions we need to cover 80% of the range of each coordinate to capture 10% of the data. (since $0.1^{1/10} \approx 0.8$), while in 2D 30% (0.316 sidelength) was suff.

Further explanation

Loosing of generalization capability

1. Image that to have K data you need 10% of the volume. How much *sidelength* (features values range) do you need?
 - From 30% to 80% moving from 2D to 10D (**loosing similarity!!!**)
2. On the other side:
 - 1D: with 0.3 *sidelength* we take 30% of data volume
 - 2D: with 0.3 *sidelength* we take 10% of data volume (the **red square**)
 - 3D: with 0.3 *sidelength* we take 3.7% of data volume (the **red cube**)
 - 10D: with 0.3 *sidelength* we take 0.0017% of data volume.

This sidelength can be not sufficient to have K data, unless we use **small K** (which can lead to overfitting)



(some) Limits of K-nn: curse of dim.

2. Low sampling density for high-dim data

- Sampling density is proportional to $l^{1/n}$ (l data, n =data dim).
 - if 100 points are sufficient to estimate a function in \mathbb{R}^1 (1-dim input),
 - 100^{10} are needed to achieve similar accuracy in \mathbb{R}^{10} (10-dim input)

(some) Limits of K-nn: curse of dim.



Dip. Informatica
University of Pisa

3. Irrelevant features: The Curse of Noisy

if the target depends on only few of many features in \mathbf{x} (e.g. 2 out of 20), we could retrieve a “similar pattern” with the similarity dominated by the large number of irrelevant features

It grows with the dimensionality

An improvement

Irrelevant features:

- We may weight features according to their relevance
 - Stretching the axes along some dimension
 - Weights can be searched by a (expensive) model selection approach or other approaches ...
- Feature selection approaches: it eliminates some variables
→ reduce input dimension

Summary:

K-nn design choices



Dip. Informatica
University of Pisa

- The metric d to measure the closeness between patterns (e.g. Euclidian, Hamming, Manhattan distance..., weights on input features).
Often this is the *key* for a successful application!
- K (number of neighbors: control underfitting/overfitting)

Often necessary to select:

- A subset of data (set of prototypes): e.g. by clustering
- A subset of features

Various approaches to deal with these issues.

Extension in ML



Dip. Informatica
University of Pisa

- Extensions to other local models in ML
 - Kernel smoothers
 - Local linear regression
 - Prototype methods
 - Case-based reasoning

K-nn in the course: some general lessons



Dip. Informatica
University of Pisa

- Too low variance is poor (rigid linear models), too high variance is dangerous (K-nn)
- Other concepts presented today that are interesting in general:
 - **Smoothing** techniques (in K-nn by increasing K)
 - **Curse of dimensionality** (the volume of the problem space increases so fast that the available data become sparse)
 - Again an **instance of the Statistical Learning Theory** bound plot: see the misclassification plot moving K
 - **Inductive Bias** for K-nn (metric based issue, which is underestimated in many books)

ML Course structure

And now?



Dip. Informatica
University of Pisa

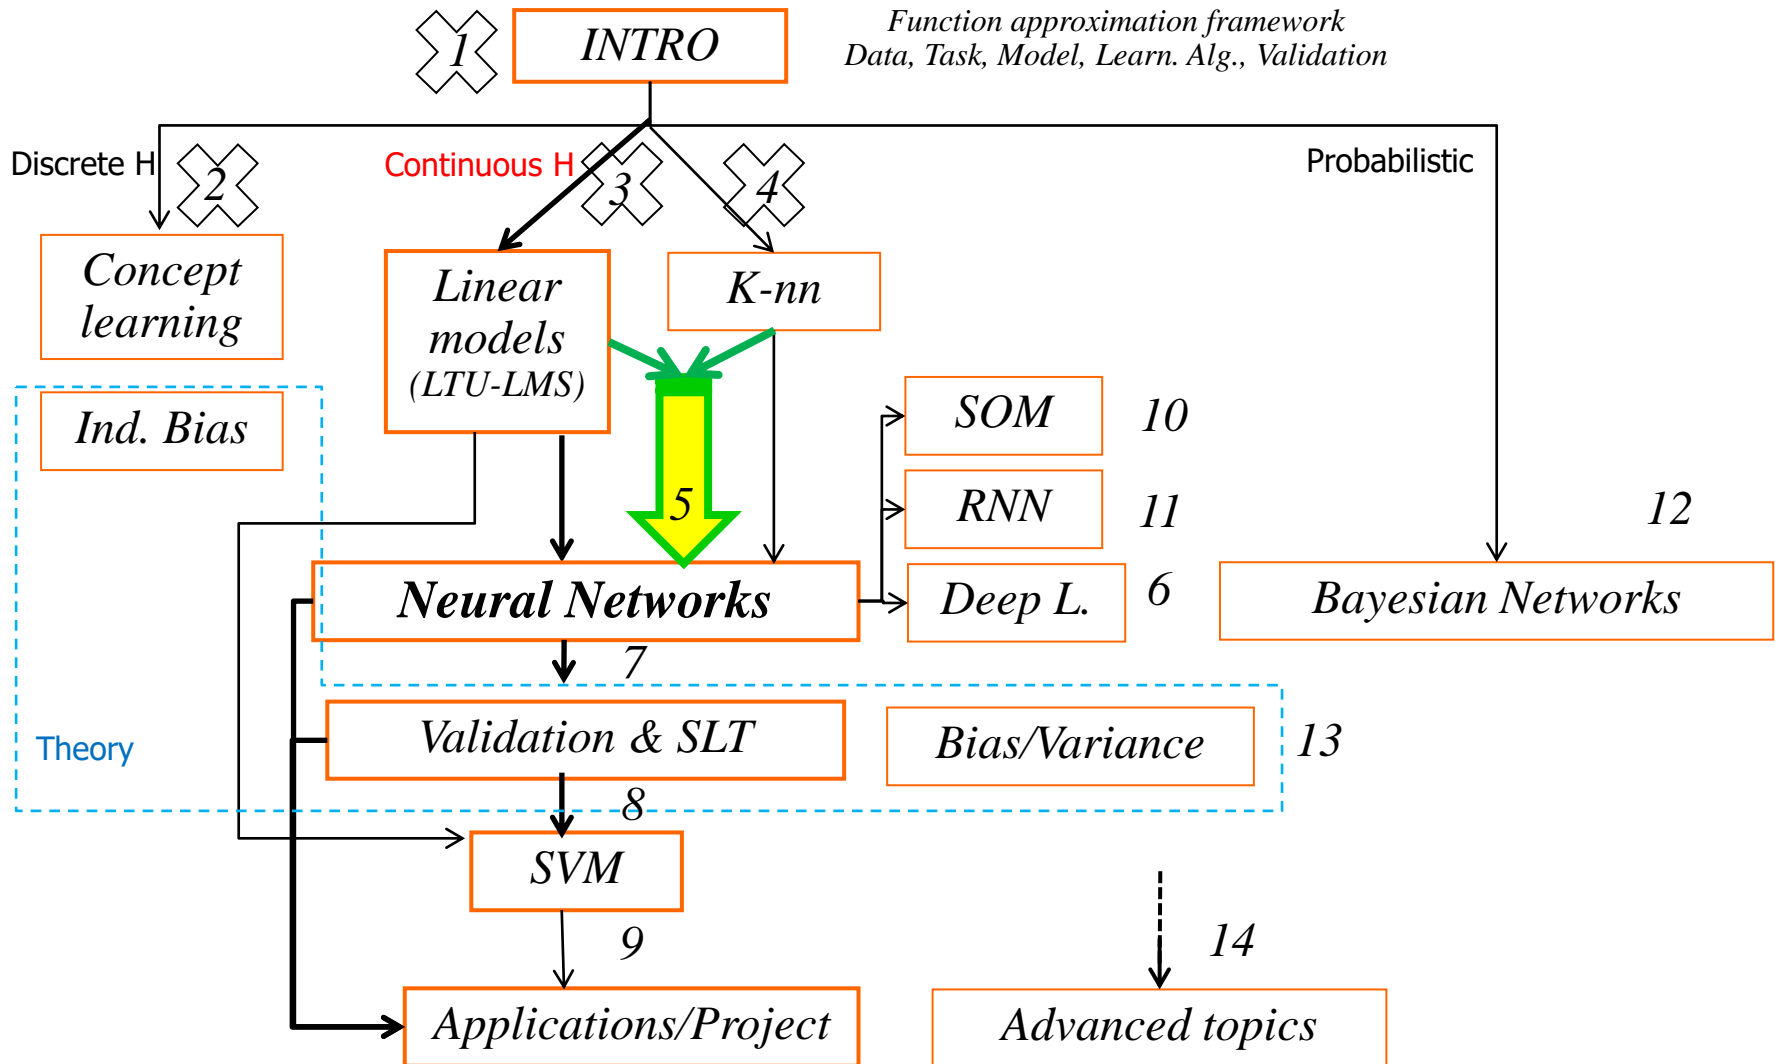
- K-nn does not build a “learned model”, not regularity/knowledge extraction/synthesis
 - It does not match the “learning” objective (that can forget the examples after the model is built)
- In the next lectures, we will look at model
 - Compact as the LTU model (all the knowledge in few parameters)
 - More flexible than the linear model (flexible as the K-nn)
 - with a suitable support to the control of the complexity

ML Course structure

Where we go



Dip. Informatica
University of Pisa



Bibliography (last 3 lectures)

- Hastie, Tibshirani, Friedman, The Elements of Statistical Learning, Springer Verlag, 2001-2017 (check the new Ed.): **chap 2**
(note: exists a on-line pdf version)
- Mitchell:
 - Linear model and LMS alg.: chap 4.4
 - K-nn: **chap 8**
- Haykin (2nd edition):
 - Linear model and LMS alg.: chap 3 up to 3.7
[details on Newton and Gauss-Newton methods are not strictly needed]
- Haykin (3rd edition):
 - Linear model and LMS alg.: chap 3
[details on Newton and Gauss-Newton and other advanced approaches are not strictly needed]

Bibliography (last 3 lectures)

- Moreover...
- Further readings (not mandatory!):
 - Gently introduction to linear models:

AIMA (Russel, Norvig, Artificial Intelligence: A Modern Approach), ed .3: **chap 18.6** (18.6.1,18.6.2,18.6.3)
 - To go in deep for **Linear least squares**

You can start from www resource as (With nice examples):
https://en.wikipedia.org/wiki/Linear_least_squares_%28mathematics%29
 - **LS in general**
http://en.wikipedia.org/wiki/Least_squares

For information

Alessio Micheli
micheli@di.unipi.it

www.di.unipi.it/groups/ciml



Dipartimento di Informatica
Università di Pisa - Italy
Micheli



**Computational Intelligence &
Machine Learning Group**